

Text moderation

WORKING WITH THE OPENAI API

James Chapman

Curriculum Manager, DataCamp

Going beyond text completions...

Completions → generate *new text output* using text prompt

Beyond completions:


- Text moderation
- Audio transcription and translation
- Combining models together

Text moderation

- Identifying inappropriate content


Traditionally,

- Moderators flag content *by-hand*
 - Time-consuming
- **Word-detection algorithms**
 - Lack *nuance* and understanding of *context*



Violation categories

- Identify violations of terms or use
- Differentiate violation type by category
 - Violence
 - Hate Speech



¹ <https://openai.com/policies/usage-policies> ² <https://platform.openai.com/docs/guides/moderation/overview>

Creating a moderations request

```
import openai

openai.api_key = "ENTER API KEY"

response = openai.Moderation.create(
    model="text-moderation-latest",
    input="I could kill for a hamburger."
)

print(response)
```

Interpreting the results

- `categories`
 - `true` / `false` indicator of category violation
- `category_scores`
 - Confidence of a violence
- `flagged`
 - `true` / `false` indicator of a violation

```
{
  "id": "modr-7G6SRC9ui7LACHMNJEWG51tvVCP1P",
  "model": "text-moderation-004",
  "results": [
    {
      "categories": {
        "hate": false,
        "hate/threatening": false,
        "self-harm": false,
        "sexual": false,
        "sexual/minors": false,
        "violence": false,
        "violence/graphic": false
      },
      "category_scores": {
        "hate": 6.755938670721662e-07,
        "hate/threatening": 1.2668371907054166e-10,
        "self-harm": 3.033803963603532e-08,
        "sexual": 3.787169688962422e-08,
        "sexual/minors": 6.083700765380229e-10,
        "violence": 5.8589016589394305e-06,
        "violence/graphic": 4.4532868059832253e-08
      },
      "flagged": false
    }
  ]
}
```

Interpreting the category scores

```
print(response["results"][0]["category_scores"])
```

```
{
  "hate": 6.755938670721662e-07,
  "hate/threatening": 1.2668371907054166e-10,
  "self-harm": 3.033803963603532e-08,
  "sexual": 3.787169688962422e-08,
  "sexual/minors": 6.083700765380229e-10,
  "violence": 5.8589016589394305e-06,
  "violence/graphic": 4.4532868059832253e-08
}
```

- Larger numbers → greater certainty of violation
- Numbers \neq probabilities

Considerations for implementing moderation

- Determine appropriate thresholds for each use case
- Stricter thresholds may result in fewer *false negatives*
- More lenient thresholds may result in fewer *false positives*

```
{  
  "hate": 6.755938670721662e-07,  
  "hate/threatening": 1.2668371907054166e-10,  
  "self-harm": 3.033803963603532e-08,  
  "sexual": 3.787169688962422e-08,  
  "sexual/minors": 6.083700765380229e-10,  
  "violence": 5.8589016589394305e-06,  
  "violence/graphic": 4.4532868059832253e-08  
}
```


Let's practice!
WORKING WITH THE OPENAI API

Speech-to-Text Transcription with Whisper

WORKING WITH THE OPENAI API

James Chapman

Curriculum Manager, DataCamp

OpenAI's Whisper

Speech-to-text capabilities:

- Transcribe audio
- Translate and transcribe audio into *English*
- Supports `mp3` , `mp4` , `mpeg` , `mpga` , `m4a` , `wav` , and `webm` (25 MB limit)
- Meeting transcripts
- Video captions



Loading audio files

Example: transcribe `meeting_recording_20230602.mp3`

```
audio_file = open("meeting_recording_20230602.mp3", "rb")
```

If the file is located in a *different directory*

```
audio_file = open("path/to/file/meeting_recording_20230602.mp3", "rb")
```

Making a request

- `Audio` endpoint

```
audio_file= open("meeting_recording_20230602.mp3", "rb")

response = openai.Audio.transcribe("whisper-1", audio_file)

print(response)
```

```
{
  "text": "Welcome everyone to the June product monthly. We'll get started in..."
}
```

The transcript

```
print(response["text"])
```

```
Welcome everyone to the June product monthly. We'll get started in just a minute.
Alright, let's get started. Today's agenda will start with a spotlight from Chris
on the new mobile user onboarding flow, then we'll review how we're tracking on
our quarterly targets, and finally, we'll finish with another spotlight from Katie
who will discuss the upcoming branding updates...
```

- Don't use sensitive or confidential recordings

Transcribing non-English languages

Afrikaans, Arabic, Armenian, Azerbaijani, Belarusian, Bosnian, Bulgarian, Catalan, Chinese, Croatian, Czech, Danish, Dutch, English, Estonian, Finnish, French, Galician, German, Greek, Hebrew, Hindi, Hungarian, Icelandic, Indonesian, Italian, Japanese, Kannada, Kazakh, Korean, Latvian, Lithuanian, Macedonian, Malay, Marathi, Maori, Nepali, Norwegian, Persian, Polish, Portuguese, Romanian, Russian, Serbian, Slovak, Slovenian, Spanish, Swahili, Swedish, Tagalog, Tamil, Thai, Turkish, Ukrainian, Urdu, Vietnamese, and Welsh.

1. `open()` audio file
2. Make a transcribe request to the `Audio` endpoint
3. Extract text from the response

Let's practice!
WORKING WITH THE OPENAI API

Speech Translation with Whisper

WORKING WITH THE OPENAI API

James Chapman

Curriculum Manager, DataCamp

Whisper's translation capabilities

- Translate and transcribe audio
- Currently limited to *English* transcripts
- Supports `mp3` , `mp4` , `mpeg` , `mpga` , `m4a` , `wav` , and `webm` (25 MB limit)



Translating audio

```
audio_file = open("non_english_audio.m4a", "rb")

response = openai.Audio.translate("whisper-1", audio_file)

print(response["text"])
```

The search volume for keywords like A I has increased rapidly since the launch of ChatGPT.

- Performance can vary wildly, depending on:
 - Audio quality
 - Audio language
 - Model's knowledge of subject matter

Bringing prompts into the mix

- Can provide `prompt` to the model (optional)
- Improve response quality by:
 - Providing an example of desired style
 - Provide context on transcript context

Example: Retaining filler words

```
prompt="Ok, ummm... this is what we should do, like, to uhhh... increase revenue."
```

Example: Provide context

```
prompt="A discussion on how to increase revenue."
```

Adding in a prompt

```
audio_file = open("non_english_audio.m4a", "rb")
prompt = "The transcript is about AI trends and ChatGPT."

response = openai.Audio.translate("whisper-1",
                                   audio_file,
                                   prompt=prompt)

print(response["text"])
```

The search volume for keywords like AI has increased rapidly since the launch of ChatGPT.

Let's practice!
WORKING WITH THE OPENAI API

Combining models


WORKING WITH THE OPENAI API

James Chapman
Curriculum Manager, DataCamp

Combining models



Combining models



- **Chaining:** Feeding output from one model into another
- Can use the *same model* multiple times
 - **Example:** validating original response
- Or *different models*
 - **Example:** summarizing meeting recordings

Example: Extracting meeting attendees

```
audio_file = open("meeting_recording_20230608.mp4", "rb")

audio_response = openai.Audio.transcribe("whisper-1", audio_file)

transcript = audio_response["text"]
prompt = "Extract the attendee names from the start of this meeting transcript: " + transcript

chat_response = openai.ChatCompletion.create(
    model="gpt-3.5-turbo",
    messages=[
        {"role": "user", "content": prompt}
    ]
)

print(chat_response['choices'][0]['message']['content'])
```

Example: Extracting meeting attendees

The meeting attendees were Otis, Paul, Elaine, Nicola, Alan, and Imran.

Note:

- **No guarantees** on model performance
- Ensure that applications are **well-tested**
- Usage should be restricted to **non-sensitive data**

Let's practice!
WORKING WITH THE OPENAI API

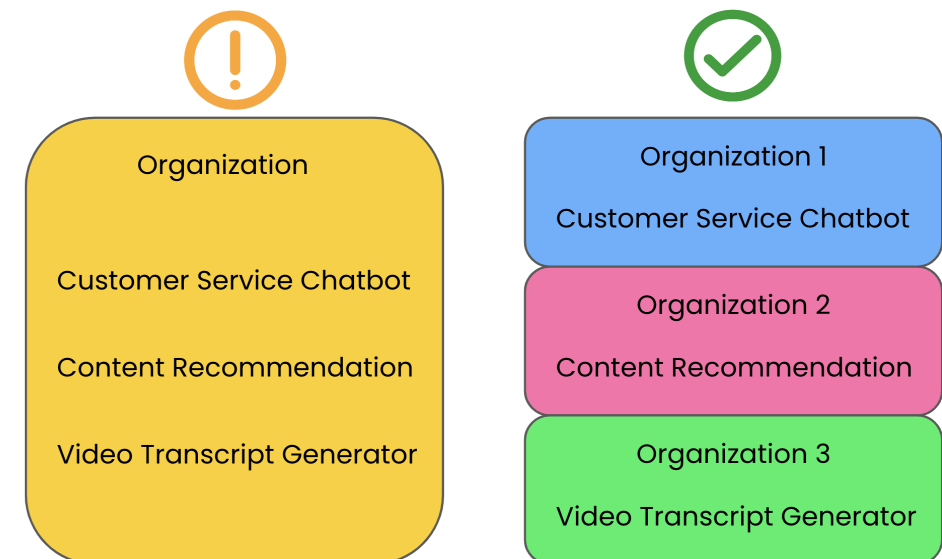

Congratulations!

WORKING WITH THE OPENAI API

James Chapman
Curriculum Manager, DataCamp



Chapter 1

- What the OpenAI API is used for
- How to create an API key
- How to create requests to the Completion endpoint
- Creating and managing organizations



Chapter 2

```
response = openai.Completion.create(  
    model="gpt-3.5-turbo-instruct",  
    prompt="When was OpenAI founded?",  
    max_tokens=20,  
    temperature=0.5  
)
```



- Solved different tasks with text completions, including:
 - Q&A
 - Text transformation
 - Content generation
 - Sentiment analysis
 - Categorization
- `max_tokens` and `temperature`
- Multi-turn conversation with `ChatCompletion`

Chapter 3

Moderation


- Detect inappropriate content

Audio

- Translate
- Transcribe

Chaining

- Automating complex tasks



Let's practice!
WORKING WITH THE OPENAI API